

2024-1

Multivariate Data Analysis

Youjin Shin
(yj.shinn@catholic.ac.kr)




Data Wrangling

◆ Data wrangling

- Data wrangling (Data manipulation) often referred to as data munging, is a critical step in the data science process. It involves cleaning, structuring, and enriching raw data into a desired format for better decision making in less time. This process is fundamental because data is rarely in a form that is immediately optimal for obtaining insights. Data wrangling makes data more usable and accessible.

Data Wrangling

***Data preparation** accounts for about 80% of the work of data scientists*



What data scientists spend the most time doing

- Building training sets: 3%
- Cleaning and organizing data: 60%
- Collecting data sets: 19%
- Mining data for patterns: 9%
- Refining algorithms: 4%
- Other: 5%

Data Wrangling

◆ Key Steps in Data Wrangling

- 1. Gathering Data:** Collecting data from various sources such as files, databases, web servers, APIs, and more.
- 2. Assessing Data:** Evaluating the data for quality and structure issues. This includes identifying and diagnosing data inconsistencies, missing values, and the need for restructuring.
- 3. Cleaning Data:** Addressing the issues identified in the assessment phase. This involves correcting, imputing, or removing erroneous data. Techniques include deduplication, filling in missing values, correcting errors, and more.



Data Wrangling

◆ Key Steps in Data Wrangling

4. **Transforming Data:** Changing the format or structure of the data to make it more suitable for analysis. This could involve normalizing data, aggregating information, pivoting tables, and more.
5. **Enriching Data:** Enhancing data by merging it with other data sources, adding derived attributes, or performing calculations that make the data more informative.
6. **Validating and Publishing Data:** Ensuring that the data meets the criteria and standards for the analytics purposes. This could involve consistency checks, quality assurance, and finally making the data available for analysis or reporting.



Lab: Data Wrangling

◆ NYC weather data

- we will using NYC daily weather data that was taken from the National Centers for Environmental Information (NCEI) API
- 1-Year data, 1 station gathers the data once a day

◆ Data meanings:

- PRCP: precipitation in millimeters
- SNOW: snowfall in millimeters
- SNWD: snow depth in millimeters
- TMAX: maximum daily temperature in Celsius
- TMIN: minimum daily temperature in Celsius
- TOBS: temperature at time of observation in Celsius
- WESF: water equivalent of snow in millimeters

◆ Some important facts to get our bearings:

- According to the National Weather Service, the coldest temperature ever recorded in Central Park was -15°F (-26.1°C) on February 9, 1934: source
- The temperature of the Sun's photosphere is approximately $5,505^{\circ}\text{C}$: source

Lab: Data Wrangling

```
# Understanding data - head()

import pandas as pd
import numpy as np

pd.set_option('display.max_columns', None)

df = pd.read_csv('data/MDA_04_dirty_data.csv')
print(df.head())
```

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
0	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
1	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
3	2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-8.3	-16.1	-12.2	NaN	False
4	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False

Lab: Data Wrangling

```
# Understanding data - head()
```

```
import pandas as pd
import numpy as np
```

```
pd.set_option('display.max_columns', None)
```

```
df = pd.read_csv('data/MDA_04_dirty_data.csv')
print(df.head())
```

<Need to check>

- Station : ?
- SNWD : -inf
- TMAX 5505.0
- Some NaN s

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
0	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
1	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
3	2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-8.3	-16.1	-12.2	NaN	False
4	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False

Lab: Data Wrangling

```
# Understanding data - describe()
```

```
print(df.describe())
```

	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF
count	765.000000	577.000000	577.0	765.000000	765.000000	398.000000	11.000000
mean	5.360392	4.202773	NaN	2649.175294	-15.914379	8.632161	16.290909
std	10.002138	25.086077	NaN	2744.156281	24.242849	9.815054	9.489832
min	0.000000	0.000000	-inf	-11.700000	-40.000000	-16.100000	1.800000
25%	0.000000	0.000000	NaN	13.300000	-40.000000	0.150000	8.600000
50%	0.000000	0.000000	NaN	32.800000	-11.100000	8.300000	19.300000
75%	5.800000	0.000000	NaN	5505.000000	6.700000	18.300000	24.900000
max	61.700000	229.000000	inf	5505.000000	23.900000	26.100000	28.700000

Lab: Data Wrangling

```
# Understanding data - info()
```

```
print(df.info())
```

<Need to check>

- 765 entries (we have 1-year data)
- # of Non-null are different
- Inclement_weather is 'object', not 'boolean'

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 765 entries, 0 to 764
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	date	765 non-null	object
1	station	765 non-null	object
2	PRCP	765 non-null	float64
3	SNOW	577 non-null	float64
4	SNWD	577 non-null	float64
5	TMAX	765 non-null	float64
6	TMIN	765 non-null	float64
7	TOBS	398 non-null	float64
8	WESF	11 non-null	float64
9	inclement_weather	408 non-null	object

```
dtypes: float64(7), object(3)
```

Lab: Data Wrangling

```
# Understanding data - info()
```

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 765 entries, 0 to 764
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	date	765 non-null	object
1	station	765 non-null	object
2	PRCP	765 non-null	float64
3	SNOW	577 non-null	float64
4	SNWD	577 non-null	float64
5	TMAX	765 non-null	float64
6	TMIN	765 non-null	float64
7	TOBS	398 non-null	float64
8	WESF	11 non-null	float64
9	inclement_weather	408 non-null	object

```
dtypes: float64(7), object(3)
```

<Need to check>

- 765 entries (we have 1-year data)
- # of Non-null are different
- Inclement_weather is 'object', not 'boolean'

Lab: Data Wrangling

Understanding Nulls : (1) WESF has only 11 non-null values: how do they look like?

```
print(df[df.WESF.isna()==False])
```

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
7	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
8	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
58	2018-01-30T00:00:00	?	1.5	13.0	inf	5505.0	-40.0	NaN	1.8	True
137	2018-03-08T00:00:00	?	28.4	NaN	NaN	5505.0	-40.0	NaN	28.7	NaN
146	2018-03-13T00:00:00	?	3.0	13.0	inf	5505.0	-40.0	NaN	3.0	True
159	2018-03-21T00:00:00	?	6.6	114.0	inf	5505.0	-40.0	NaN	8.6	True
162	2018-03-21T00:00:00	?	6.6	114.0	inf	5505.0	-40.0	NaN	8.6	True
186	2018-04-02T00:00:00	?	14.0	152.0	inf	5505.0	-40.0	NaN	15.2	True
678	2018-11-16T00:00:00	?	47.0	152.0	inf	5505.0	-40.0	NaN	24.9	True
679	2018-11-16T00:00:00	?	47.0	152.0	inf	5505.0	-40.0	NaN	24.9	True
680	2018-11-16T00:00:00	?	47.0	152.0	inf	5505.0	-40.0	NaN	24.9	True

Lab: Data Wrangling

Understanding Nulls : (2) How many entries have at least one null value?

```
contain_nulls = df[
    df.SNOW.isna() | df.SNWD.isna() | df.TOBS.isna()
    | df.WESF.isna() | df.inclement_weather.isna()
]
print(contain_nulls.shape[0])
print(contain_nulls.head(10))
```

765

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
0	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
1	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
3	2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-8.3	-16.1	-12.2	NaN	False
4	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
5	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
6	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
7	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
8	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
9	2018-01-05T00:00:00	?	0.3	NaN	NaN	5505.0	-40.0	NaN	NaN	NaN

Lab: Data Wrangling

Understanding inf/-inf: (1) Find the number of inf/-inf values per column in the df

```
def get_inf_count(df):  
    return {  
        col: df[df[col].isin([np.inf, -np.inf])].shape[0] for col in df.columns  
    }  
  
print(get_inf_count(df))
```

```
{'date': 0, 'station': 0, 'PRCP': 0, 'SNOW': 0, 'SNWD': 577, 'TMAX': 0, 'TMIN': 0, 'TOBS': 0, 'WESF': 0,  
'inclement_weather': 0}
```

- (CF) # of null values per column?
 - ➔ We already get information of Nulls from info()
 - ➔ 765-(non-null)

Lab: Data Wrangling

```
# Understanding inf/-inf: (2) Take a deep look at the 'SNWD(Snow depth)'
```

```
# Since snow depth (SNWD) and snowfall amount (SNOW) are closely related, let's interpret SNWD's inf/-inf information through SNOW.
```

```
snow = pd.DataFrame({  
    'np.inf Snow Depth': df[df.SNWD == np.inf].SNOW.describe(),  
    '-np.inf Snow Depth': df[df.SNWD == -np.inf].SNOW.describe()  
})  
  
print(snow)
```

	np.inf Snow Depth	-np.inf Snow Depth
count	24.000000	553.0
mean	101.041667	0.0
std	74.498018	0.0
min	13.000000	0.0
25%	25.000000	0.0
50%	120.500000	0.0
75%	152.000000	0.0
max	229.000000	0.0


Lab: Data Wrangling

```
# Understanding inf/-inf: (2) Take a deep look at the 'SNWD(Snow depth)'
```

```
# Since snow depth (SNWD) and snowfall amount (SNOW) are closely related, let's interpret SNWD's inf/-inf information through SNOW.
```

```
snow = pd.DataFrame({  
    'np.inf Snow Depth': df[df.SNWD == np.inf].SNOW.describe(),  
    '-np.inf Snow Depth': df[df.SNWD == -np.inf].SNOW.describe()  
})  
  
print(snow)
```

	np.inf Snow Depth	-np.inf Snow Depth
count	24.000000	553.0
mean	101.041667	0.0
std	74.498018	0.0
min	13.000000	0.0
25%	25.000000	0.0
50%	120.500000	0.0
75%	152.000000	0.0
max	229.000000	0.0



-inf -> 0 -> NO SNOW

Lab: Data Wrangling

```
# Understanding 'data' and 'station'
```

```
print(df.describe(include='object'))
```

	date	station	inclement_weather
count	765	765	408
unique	324	2	2
top	2018-07-05T00:00:00	GHCND:USC00280907	False
freq	8	398	384

- 'data'

- We have 1-year data -> 765 rows : Need to check
- 'data' has 324 unique rows (Some might be duplicated(765:365) and missing(365:324))
- A specific date appears up to 8 times (2018-07-05T00:00:00)

- 'station'

- We have 1 station, but # of unique data is 2 -> 'GHCND:USC00280907' and '?'
- 'GHCND:USC00280907' appears 398 times -> '?' appears (756-398) = 367 times

Lab: Data Wrangling

```
# Understanding 'data' and 'station' : Duplication information
print(df[df.duplicated(keep=False)].head(10))
```

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
0	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
1	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
4	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
5	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
6	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
7	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
8	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
9	2018-01-05T00:00:00	?	0.3	NaN	NaN	5505.0	-40.0	NaN	NaN	NaN
10	2018-01-05T00:00:00	?	0.3	NaN	NaN	5505.0	-40.0	NaN	NaN	NaN
20	2018-01-12T00:00:00	GHCND:USC00280907	1.3	0.0	-inf	9.4	0.6	7.8	NaN	False
21	2018-01-12T00:00:00	?	0.5	NaN	NaN	5505.0	-40.0	NaN	NaN	NaN
22	2018-01-12T00:00:00	?	0.5	NaN	NaN	5505.0	-40.0	NaN	NaN	NaN
23	2018-01-12T00:00:00	GHCND:USC00280907	1.3	0.0	-inf	9.4	0.6	7.8	NaN	False
24	2018-01-12T00:00:00	GHCND:USC00280907	1.3	0.0	-inf	9.4	0.6	7.8	NaN	False
26	2018-01-14T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	2.2	-11.1	-11.1	NaN	False
27	2018-01-14T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
28	2018-01-14T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
29	2018-01-14T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	2.2	-11.1	-11.1	NaN	False
31	2018-01-16T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-1.1	-9.4	-4.4	NaN	False

Same date
Diff. stations

'?' station
has lots NaN
values,
and 5505s

Lab: Data Wrangling

- ◆ Handling the data!!
 - ◆ Remove duplications : Combine the row having same date into one row.
- BEFORE THAT**, Let's consider the values that we should not erase carelessly!
- We will Keep the WESF values first

- ◆ QUIZ: Find unique 'station' values where the WSF value is not Null

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
0	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
1	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2	2018-01-01T00:00:00	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
3	2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-8.3	-16.1	-12.2	NaN	False
4	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN	False
X	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	13	False
X	2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-inf	-4.4	-13.9	-13.3	13	False
X	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
X	2018-01-04T00:00:00	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
9	2018-01-05T00:00:00	?	0.3	NaN	NaN	5505.0	-40.0	NaN	NaN	NaN

For example:

GHCND:USC00280907, ?

Lab: Data Wrangling

- ◆ QUIZ: Find unique 'station' values where the WSF value is not Null

```
# Understanding 'data' and 'station' : Unique value of the station whose WESF value is not null  
  
print(df[df.WESF.notna()].station.unique())
```

```
['?']
```

Lab: Data Wrangling

- ◆ QUIZ: Find unique 'station' values where the WESF value is not Null

```
# Understanding 'data' and 'station' : Unique value of the station whose WESF value is not null  
  
print(df[df.WESF.notna()].station.unique())
```

```
['?']
```

- we only have values for the WESF column when the station is '?'

Lab: Data Wrangling

◆ Handling duplications : Combine the row having same date into one row.

- We will drop rows based on date.
- Keep the WESF values first. (index = date)
- If there are two types of station information for the same date, remove '?'
- Delete the 'station' column as it is no longer needed.
- Combine the WESF values based on date.

```
# 1. make the date a datetime (We don't need time (hour/minutes...) information.)
```

```
# 2. save the WEST information
```

```
# 3. sort '?' to the bottom
```

```
# 4. drop duplicates based on the date column keeping the first occurrence
```

```
# 5. remove the station column because we are done with it
```

```
# 6. Insert the saved WESF value.
```

Lab: Data Wrangling

```
# 1. make the date a datetime (We don't need time (hour/minutes...) information.)
```

```
df.date = pd.to_datetime(df.date)
```

```
# 2. save the WEST information
```

```
station_qm_wesf = df[df.station == '?'].drop_duplicates('date').set_index('date').WESF
```

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	...
0	2018-01-01	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	
1	2018-01-01	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	
2	2018-01-01	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	
7	2018-01-04	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	
8	2018-01-04	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	
9	2018-01-05	?	0.3	NaN	NaN	5505.0	-40.0	NaN	NaN	
10	2018-01-05	?	0.3	NaN	NaN	5505.0	-40.0	NaN	NaN	
21	2018-01-12	?	0.5	NaN	NaN	5505.0	-40.0	NaN	NaN	
22	2018-01-12	?	0.5	NaN	NaN	5505.0	-40.0	NaN	NaN	
27	2018-01-14	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	
28	2018-01-14	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	

```
df[df.station == '?']
```



	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	...
0	2018-01-01	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	
7	2018-01-04	?	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	
9	2018-01-05	?	0.3	NaN	NaN	5505.0	-40.0	NaN	NaN	
21	2018-01-12	?	0.5	NaN	NaN	5505.0	-40.0	NaN	NaN	
27	2018-01-14	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	

```
drop_duplicates('date')
```



```
date
2018-01-01    NaN
2018-01-04    19.3
2018-01-05    NaN
2018-01-12    NaN
2018-01-14    NaN
Name: WESF, dtype: float64
```

```
set_index('date').WESF
```

Lab: Data Wrangling

```
# 3. sort '?' to the bottom
df.sort_values('station', ascending=False, inplace=True)

# 4. drop duplicates based on the date column keeping the first occurrence
# which will be the valid station if it has data
df_deduped = df.drop_duplicates('date')
```

	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	...
80	2018-02-10	GHCND:USC00280907	0.0	0.0	-inf	0.0	-8.3	0.0	NaN	
122	2018-03-01	GHCND:USC00280907	0.0	0.0	-inf	14.4	0.0	3.9	NaN	
120	2018-03-01	GHCND:USC00280907	0.0	0.0	-inf	14.4	0.0	3.9	NaN	
116	2018-02-27	GHCND:USC00280907	0.0	0.0	-inf	10.6	0.0	0.0	NaN	
115	2018-02-27	GHCND:USC00280907	0.0	0.0	-inf	10.6	0.0	0.0	NaN	
..	
54	2018-01-26	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	
53	2018-01-26	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	
99	2018-02-16	?	3.3	NaN	NaN	5505.0	-40.0	NaN	NaN	
118	2018-02-28	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	
0	2018-01-01	?	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	



	date	station	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	\
80	2018-02-10	GHCND:USC00280907	0.0	0.0	-inf	0.0	-8.3	0.0	
122	2018-03-01	GHCND:USC00280907	0.0	0.0	-inf	14.4	0.0	3.9	
116	2018-02-27	GHCND:USC00280907	0.0	0.0	-inf	10.6	0.0	0.0	
..	
54	2018-01-26	?	0.0	0.0	-inf	5505.0	-40.0	NaN	

Lab: Data Wrangling

```
# 5. remove the station column because we are done with it
df_deduped = df_deduped.drop(columns='station').set_index('date').sort_index()

# 6. Insert the saved WESF value.
# combine_first(): Combine two DataFrame objects by filling null values in one DataFrame with non-null values

df_deduped = df_deduped.assign(
    WESF=lambda x: x.WESF.combine_first(station_qm_wesf)
)

print(df_deduped.shape)
```

	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
date								
...								
2018-03-19	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2018-03-21	0.0	0.0	-inf	2.8	-2.8	0.6	NaN	False
2018-03-22	17.3	178.0	inf	1.7	-1.7	0.0	NaN	True
2018-03-23	0.0	0.0	-inf	8.9	0.0	1.1	NaN	False
2018-03-25	0.0	0.0	-inf	8.3	-1.1	-0.6	NaN	False
...								

(324, 8)

	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
date								
...								
2018-03-19	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN	NaN
2018-03-21	0.0	0.0	-inf	2.8	-2.8	0.6	8.6	False
2018-03-22	17.3	178.0	inf	1.7	-1.7	0.0	NaN	True
2018-03-23	0.0	0.0	-inf	8.9	0.0	1.1	NaN	False
2018-03-25	0.0	0.0	-inf	8.3	-1.1	-0.6	NaN	False
...								

Lab: Data Wrangling

◆ Handling nulls - Remove

- ① Remove if at least one null exist

```
print(df_deduped.dropna().shape)  
Print(df_deduped)
```

Do not apply

(4, 8)

	PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
date								
2018-01-30	0.0	0.0	-inf	6.7	-1.7	-0.6	1.8	False
2018-03-13	4.1	51.0	inf	5.6	-3.9	0.0	3.0	True
2018-03-21	0.0	0.0	-inf	2.8	-2.8	0.6	8.6	False
2018-04-02	9.1	127.0	inf	12.8	-1.1	-1.1	15.2	True

Lab: Data Wrangling

◆ Handling nulls - Remove

- ② Remove if there are nulls in certain columns (all)

Do not apply

```
df_deduped = df_deduped.dropna(  
    how='all', subset=['inclement_weather', 'SNOW', 'SNWD']  
)  
print(df_deduped)
```

PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
date							
2018-01-01	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN
2018-01-02	0.0	0.0	-inf	-8.3	-16.1	-12.2	NaN
2018-01-03	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN
2018-01-04	20.6	229.0	inf	5505.0	-40.0	NaN	19.3
2018-01-05	14.2	127.0	inf	-4.4	-13.9	-13.9	NaN
2018-01-06	0.0	0.0	-inf	-10.0	-15.6	-15.0	NaN
2018-01-07	0.0	0.0	-inf	-11.7	-17.2	-16.1	NaN
2018-01-08	0.0	0.0	-inf	-7.8	-16.7	-8.3	NaN
2018-01-10	0.0	0.0	-inf	5.0	-7.8	-7.8	NaN
2018-01-11	0.0	0.0	-inf	4.4	-7.8	1.1	NaN
2018-01-12	1.3	0.0	-inf	9.4	0.6	7.8	NaN
2018-01-13	17.5	NaN	NaN	5505.0	-40.0	NaN	NaN
2018-01-14	0.0	0.0	-inf	2.2	-11.1	-11.1	NaN
2018-01-15	0.0	0.0	-inf	-5.0	-11.7	-8.9	NaN
2018-01-16	0.0	0.0	-inf	-1.1	-9.4	-4.4	NaN



PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather
date							
2018-01-01	0.0	0.0	-inf	5505.0	-40.0	NaN	NaN
2018-01-02	0.0	0.0	-inf	-8.3	-16.1	-12.2	NaN
2018-01-03	0.0	0.0	-inf	-4.4	-13.9	-13.3	NaN
2018-01-04	20.6	229.0	inf	5505.0	-40.0	NaN	19.3
2018-01-05	14.2	127.0	inf	-4.4	-13.9	-13.9	NaN
2018-01-06	0.0	0.0	-inf	-10.0	-15.6	-15.0	NaN
2018-01-07	0.0	0.0	-inf	-11.7	-17.2	-16.1	NaN
2018-01-08	0.0	0.0	-inf	-7.8	-16.7	-8.3	NaN
2018-01-10	0.0	0.0	-inf	5.0	-7.8	-7.8	NaN
2018-01-11	0.0	0.0	-inf	4.4	-7.8	1.1	NaN
2018-01-12	1.3	0.0	-inf	9.4	0.6	7.8	NaN
2018-01-14	0.0	0.0	-inf	2.2	-11.1	-11.1	NaN
2018-01-15	0.0	0.0	-inf	-5.0	-11.7	-8.9	NaN
2018-01-16	0.0	0.0	-inf	-1.1	-9.4	-4.4	NaN

Lab: Data Wrangling

◆ Handling nulls - Remove

- ③ Remove only if there are nulls in all columns

```
df_deduped = df_deduped.dropna(how='all')  
print(df_deduped.shape)
```

```
(324, 8)
```

Lab: Data Wrangling

◆ Handling nulls - Fill

- ① Filling nulls with a constant value

```
print(df_deduped['WESF'].fillna(0, inplace=True))  
print(df_deduped)
```

PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	date
2018-01-01	0.0	0.0	-inf	5505.0	-40.0	NaN	0.0	NaN
2018-01-02	0.0	0.0	-inf	-8.3	-16.1	-12.2	0.0	False
2018-01-03	0.0	0.0	-inf	-4.4	-13.9	-13.3	0.0	False
2018-01-04	20.6	229.0	inf	5505.0	-40.0	NaN	19.3	True
2018-01-05	14.2	127.0	inf	-4.4	-13.9	-13.9	0.0	True

```
df_deduped = df_deduped.assign(  
    TMAX=lambda x: x.TMAX.replace(5505, np.nan),  
    TMIN=lambda x: x.TMIN.replace(-40, np.nan),  
)  
print(df_deduped)
```

PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	date
2018-01-01	0.0	0.0	-inf	NaN	NaN	NaN	0.0	NaN
2018-01-02	0.0	0.0	-inf	-8.3	-16.1	-12.2	0.0	False
2018-01-03	0.0	0.0	-inf	-4.4	-13.9	-13.3	0.0	False
2018-01-04	20.6	229.0	inf	NaN	NaN	NaN	19.3	True
2018-01-05	14.2	127.0	inf	-4.4	-13.9	-13.9	0.0	True

Lab: Data Wrangling

◆ Handling nulls - Fill

- ② Filling nulls with the preceding (forward fill) or succeeding (backward fill) values.

```
df_deduped=df_deduped.assign(  
    TMAX=lambda x: x.TMAX.fillna(method='ffill'),  
    TMIN=lambda x: x.TMIN.fillna(method='ffill')  
)  
  
print(df_deduped)
```

#backward fill : 'bfill'

PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	
2018-01-01	0.0	0.0	-inf	NaN	NaN	NaN	0.0	NaN
2018-01-02	0.0	0.0	-inf	-8.3	-16.1	-12.2	0.0	False
2018-01-03	0.0	0.0	-inf	-4.4	-13.9	-13.3	0.0	False
2018-01-04	20.6	229.0	inf	NaN	NaN	NaN	19.3	True
2018-01-05	14.2	127.0	inf	-4.4	-13.9	-13.9	0.0	True



PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	
2018-01-01	0.0	0.0	-inf	NaN	NaN	NaN	0.0	NaN
2018-01-02	0.0	0.0	-inf	-8.3	-16.1	-12.2	0.0	False
2018-01-03	0.0	0.0	-inf	-4.4	-13.9	-13.3	0.0	False
2018-01-04	20.6	229.0	inf	-4.4	-13.9	NaN	19.3	True
2018-01-05	14.2	127.0	inf	-4.4	-13.9	-13.9	0.0	True

Lab: Data Wrangling

◆ Handling nulls - Fill

- ③ Filling nulls with the clip() method to cap values at a specific minimum and/or maximum threshold.

```
df_deduped=df_deduped.assign(  
    SNWD=lambda x: x.SNWD.clip(0, x.SNOW)  
)  
  
print(df_deduped)
```

PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	
2018-01-01	0.0	0.0	-inf	NaN	NaN	NaN	0.0	NaN
2018-01-02	0.0	0.0	-inf	-8.3	-16.1	-12.2	0.0	False
2018-01-03	0.0	0.0	-inf	-4.4	-13.9	-13.3	0.0	False
2018-01-04	20.6	229.0	inf	NaN	NaN	NaN	19.3	True
2018-01-05	14.2	127.0	inf	-4.4	-13.9	-13.9	0.0	True



PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	
2018-01-01	0.0	0.0	0.0	NaN	NaN	NaN	0.0	NaN
2018-01-02	0.0	0.0	0.0	-8.3	-16.1	-12.2	0.0	False
2018-01-03	0.0	0.0	0.0	-4.4	-13.9	-13.3	0.0	False
2018-01-04	20.6	229.0	229.0	-4.4	-13.9	NaN	19.3	True
2018-01-05	14.2	127.0	127.0	-4.4	-13.9	-13.9	0.0	True

Lab: Data Wrangling

◆ Handling nulls - Impute

- ① Use fillna() with other types of calculations.
- We replace missing values of TMAX with the median of all TMAX values, TMIN with the median of all TMIN values, and TOBS to the average of the TMAX and TMIN values. Since we place TOBS last, we have access to the imputed values for TMIN and TMAX in the calculation

```
df_deduped=df_deduped.assign(  
    TMAX=lambda x: x.TMAX.fillna(x.TMAX.median()),  
    TMIN=lambda x: x.TMIN.fillna(x.TMIN.median()),  
    # average of TMAX and TMIN  
    TOBS=lambda x: x.TOBS.fillna((x.TMAX + x.TMIN) / 2)  
)  
print(df_deduped)
```

PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	
2018-01-01	0.0	0.0	0.0	NaN	NaN	NaN	0.0	NaN
2018-01-02	0.0	0.0	0.0	-8.3	-16.1	-12.2	0.0	False
2018-01-03	0.0	0.0	0.0	-4.4	-13.9	-13.3	0.0	False
2018-01-04	20.6	229.0	229.0	-4.4	-13.9	NaN	19.3	True
2018-01-05	14.2	127.0	127.0	-4.4	-13.9	-13.9	0.0	True



PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	
2018-01-01	0.0	0.0	0.0	16.1	6.7	11.40	0.0	NaN
2018-01-02	0.0	0.0	0.0	-8.3	-16.1	-12.20	0.0	False
2018-01-03	0.0	0.0	0.0	-4.4	-13.9	-13.30	0.0	False
2018-01-04	20.6	229.0	229.0	-4.4	-13.9	-9.15	19.3	True
2018-01-05	14.2	127.0	127.0	-4.4	-13.9	-13.90	0.0	True

Lab: Data Wrangling

◆ Handling nulls - Impute

② Interpolate with the interpolate() method

- We specify the method parameter with the interpolation strategy to use. There are many options, but we will stick with the default of 'linear', which will treat values as evenly spaced and place missing values in the middle of existing ones. We have some missing data, so we will reindex first. Look at January 9th, which we didn't have before—the values for TMAX, TMIN, and TOBS are the average of values the day prior (January 8th) and the day after (January 10th):

```
print(df_deduped.head(10))
df_deduped=df_deduped.reindex(
    pd.date_range('2018-01-01', '2018-12-31', freq='D')).apply(lambda x: x.interpolate())
print(df_deduped.head(10))
```

PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	
date								
2018-01-01	0.0	0.0	0.0	16.1	6.7	11.40	0.0	NaN
2018-01-02	0.0	0.0	0.0	-8.3	-16.1	-12.20	0.0	False
2018-01-03	0.0	0.0	0.0	-4.4	-13.9	-13.30	0.0	False
2018-01-04	20.6	229.0	229.0	-4.4	-13.9	-9.15	19.3	True
2018-01-05	14.2	127.0	127.0	-4.4	-13.9	-13.90	0.0	True
2018-01-06	0.0	0.0	0.0	-10.0	-15.6	-15.00	0.0	False
2018-01-07	0.0	0.0	0.0	-11.7	-17.2	-16.10	0.0	False
2018-01-08	0.0	0.0	0.0	-7.8	-16.7	-8.30	0.0	False
2018-01-10	0.0	0.0	0.0	5.0	-7.8	-7.80	0.0	False
2018-01-11	0.0	0.0	0.0	4.4	-7.8	1.10	0.0	False



PRCP	SNOW	SNWD	TMAX	TMIN	TOBS	WESF	inclement_weather	
2018-01-01	0.0	0.0	0.0	16.1	6.70	11.40	0.0	NaN
2018-01-02	0.0	0.0	0.0	-8.3	-16.10	-12.20	0.0	False
2018-01-03	0.0	0.0	0.0	-4.4	-13.90	-13.30	0.0	False
2018-01-04	20.6	229.0	229.0	-4.4	-13.90	-9.15	19.3	True
2018-01-05	14.2	127.0	127.0	-4.4	-13.90	-13.90	0.0	True
2018-01-06	0.0	0.0	0.0	-10.0	-15.60	-15.00	0.0	False
2018-01-07	0.0	0.0	0.0	-11.7	-17.20	-16.10	0.0	False
2018-01-08	0.0	0.0	0.0	-7.8	-16.70	-8.30	0.0	False
2018-01-09	0.0	0.0	0.0	-1.4	-12.25	-8.05	0.0	NaN
2018-01-10	0.0	0.0	0.0	5.0	-7.80	-7.80	0.0	False

*Default: interpolate missing values linearly

